1. Setup

1. Setup

Loss: mean squared error (MSE) of the future count of each bike type.

1. Setup

Loss: mean squared error (MSE) of the future count of each bike type.

Score: skill = 1 − MSE_model / MSE_persistence. Persistence = doing nothing (count now carried forward). 0 = no better than doing nothing, 1 = perfect.

1. Setup

Loss: mean squared error (MSE) of the future count of each bike type.

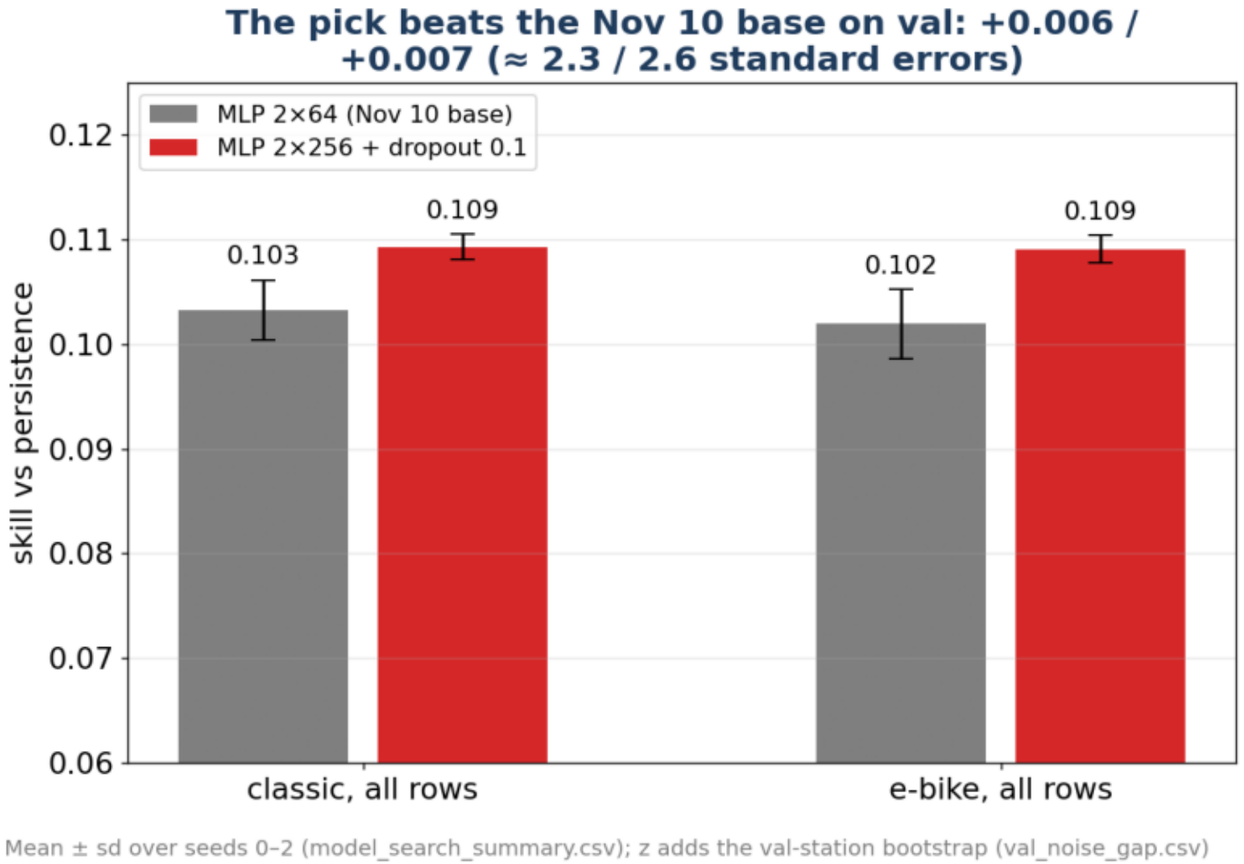
Score: skill = 1 − MSE_model / MSE_persistence. Persistence = doing nothing (count now carried forward). 0 = no better than doing nothing, 1 = perfect.

We pick models by their skill on the val days (Oct 19–23). The test days were scored once, at the end.

2. What we tested, and what we conclude

2. What we tested, and what we conclude

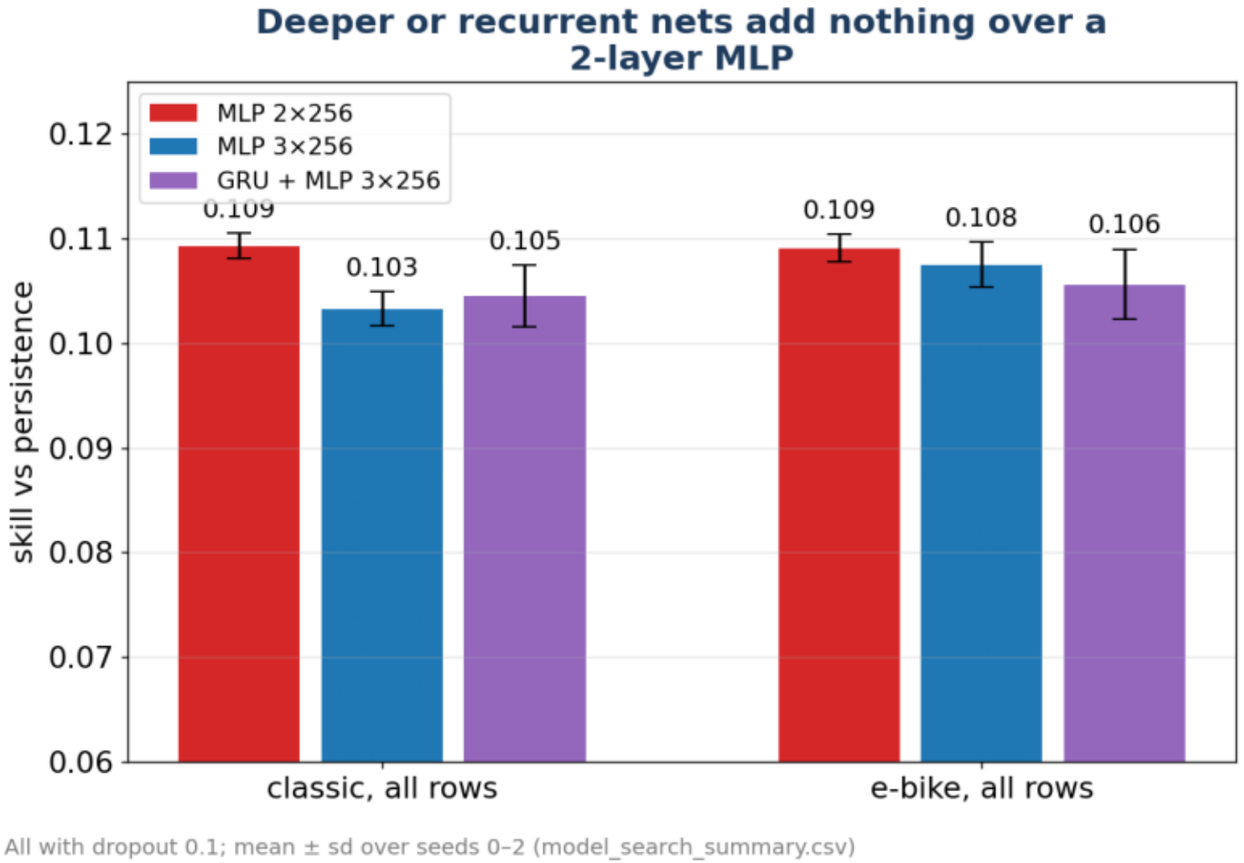
Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).



2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

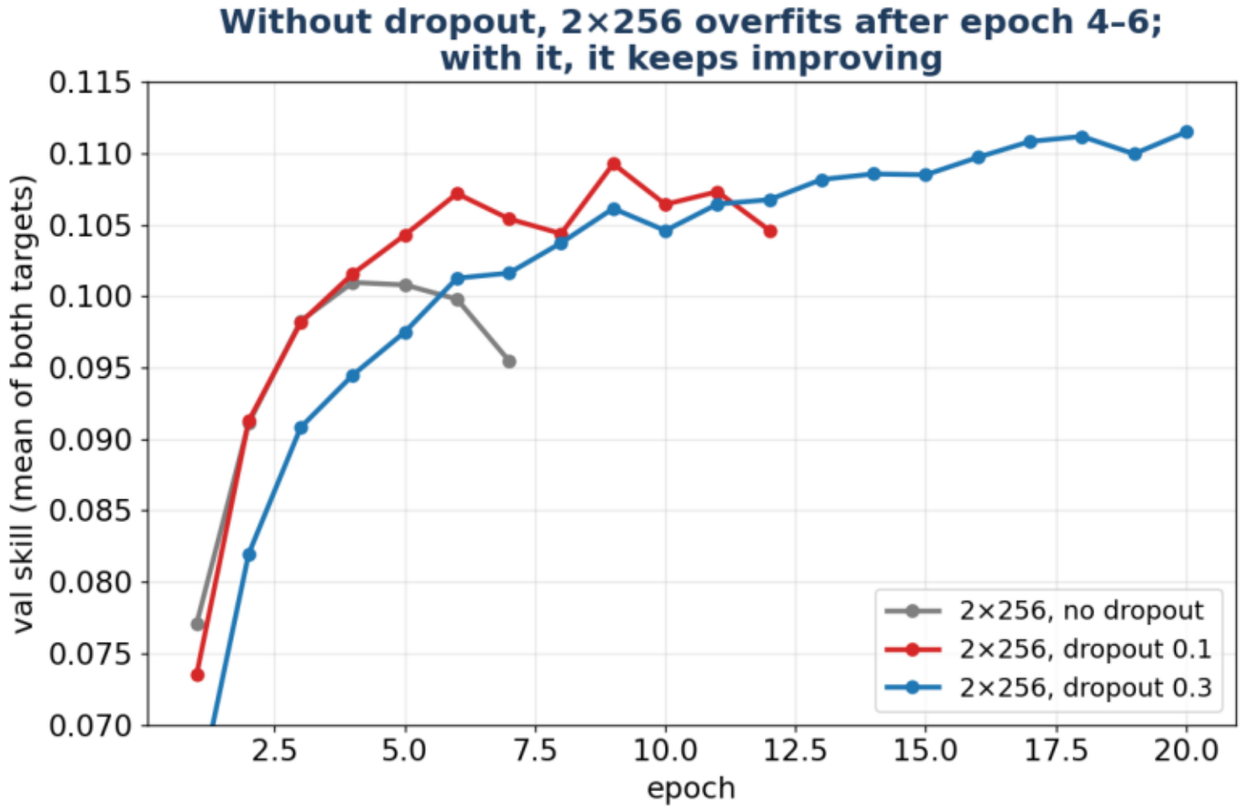


2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

Dropout is the main gain: without it, the bigger nets overfit after epoch 4–6.



Mean over seeds 0–2, epochs all three seeds reached (model_search_curves.csv)

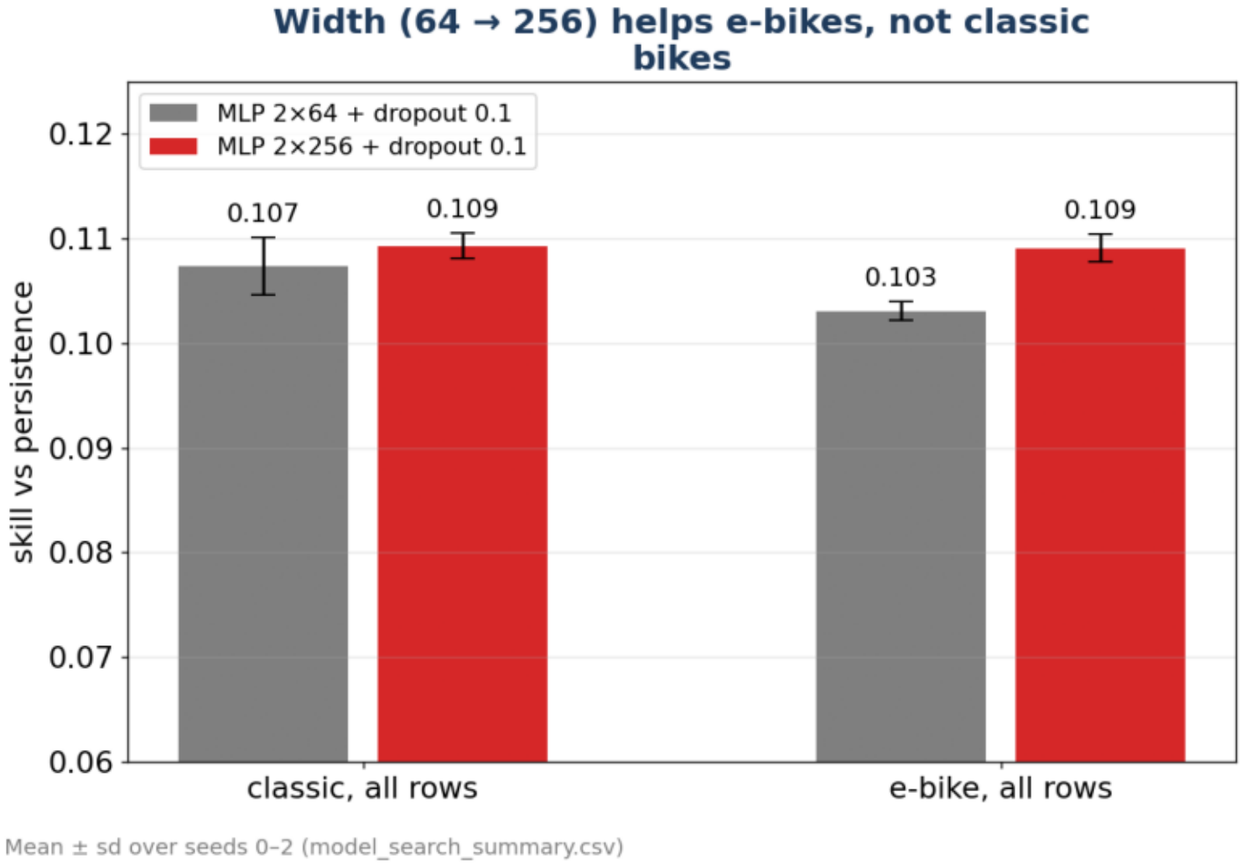
2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

Dropout is the main gain: without it, the bigger nets overfit after epoch 4–6.

Width (64 → 256) helps e-bikes only.



2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

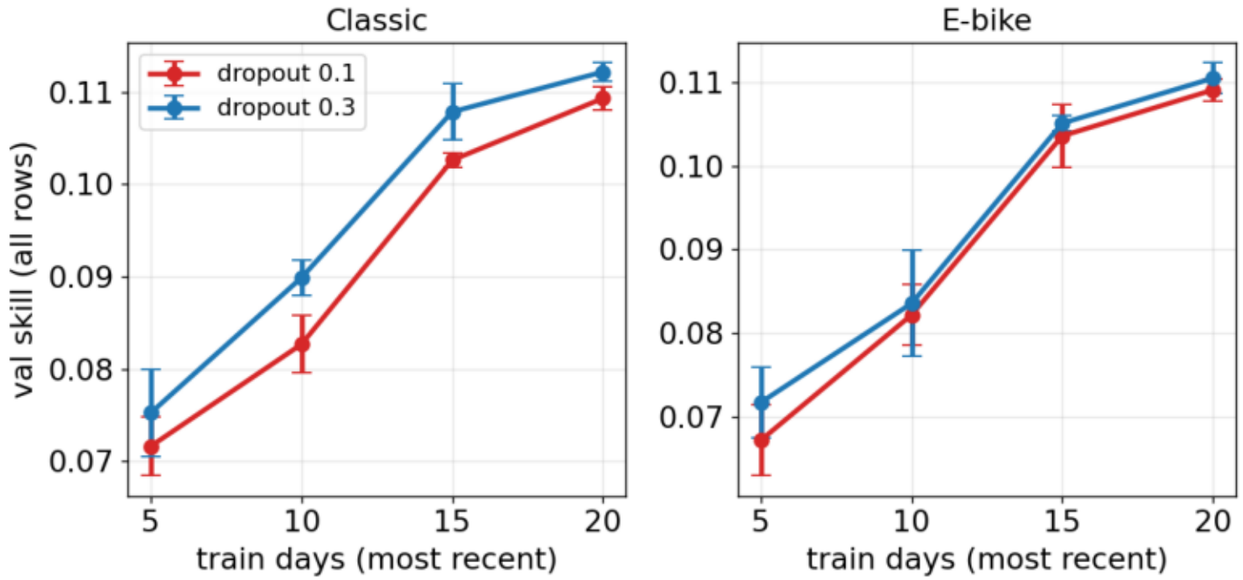
Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

Dropout is the main gain: without it, the bigger nets overfit after epoch 4–6.

Width (64 → 256) helps e-bikes only.

Why not dropout 0.3? Small lead (+0.003 / +0.001), nil on low rows, 2× the epochs, and it does not grow with data.

Dropout 0.3 leads 0.1 by a little, and the lead does not grow with data (classic +0.004 / +0.007 / +0.005 / +0.003)



Mean ± sd over seeds 0–2; same val days (data_size_summary.csv, 44_data_size.py)

2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

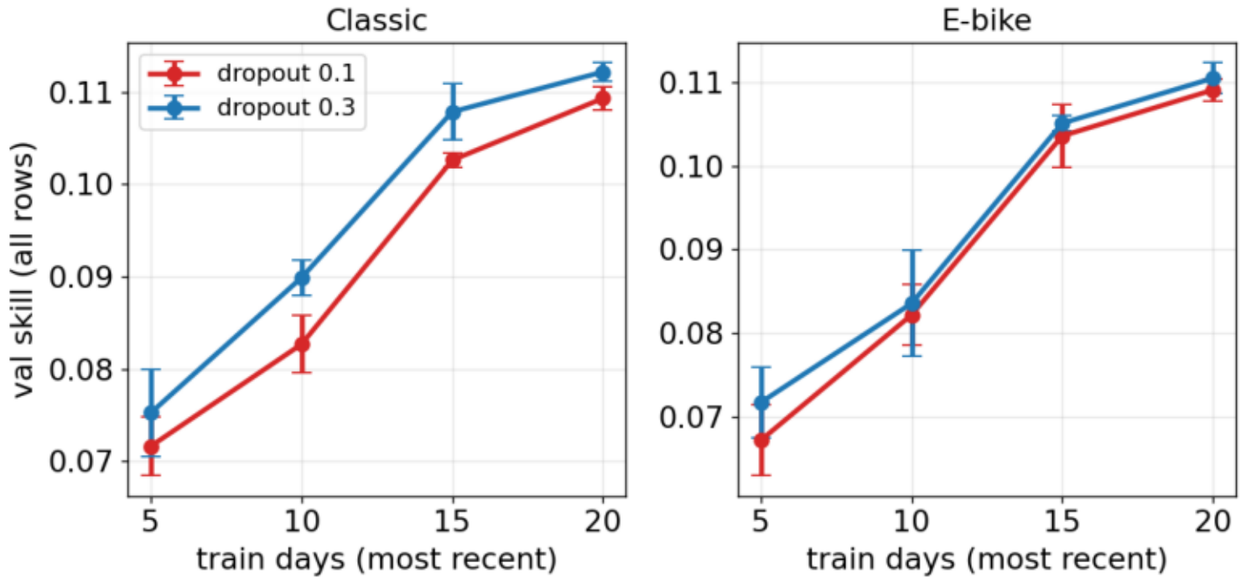
Dropout is the main gain: without it, the bigger nets overfit after epoch 4–6.

Width (64 → 256) helps e-bikes only.

Why not dropout 0.3? Small lead (+0.003 / +0.001), nil on low rows, 2× the epochs, and it does not grow with data.

More data helps most: about +0.009 / +0.008 per extra week of training days.

More train days still help: 15 → 20 days adds +0.007 / +0.006 (≈ +0.009 / +0.008 per week), but the gain is slowing



Mean ± sd over seeds 0–2; same val days (data_size_summary.csv, 44_data_size.py)

2. What we tested, and what we conclude

Pick: MLP 2×256 + dropout 0.1: +0.006 / +0.007 over the Nov 10 base (≈ 2.3–2.6 SE).

Architecture: a plain MLP. 3 layers or a GRU add nothing; no LSTM, as there are only 3 lag steps.

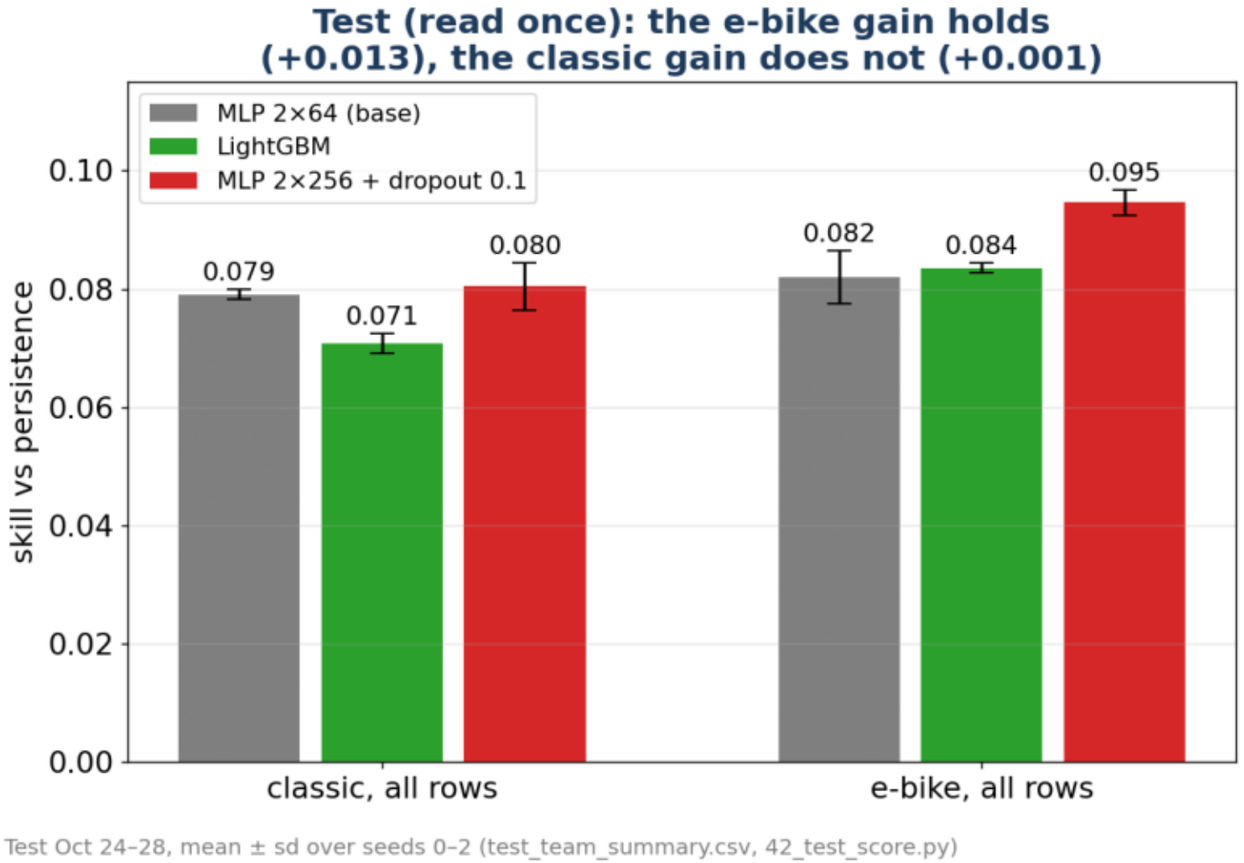
Dropout is the main gain: without it, the bigger nets overfit after epoch 4–6.

Width (64 → 256) helps e-bikes only.

Why not dropout 0.3? Small lead (+0.003 / +0.001), nil on low rows, 2× the epochs, and it does not grow with data.

More data helps most: about +0.009 / +0.008 per extra week of training days.

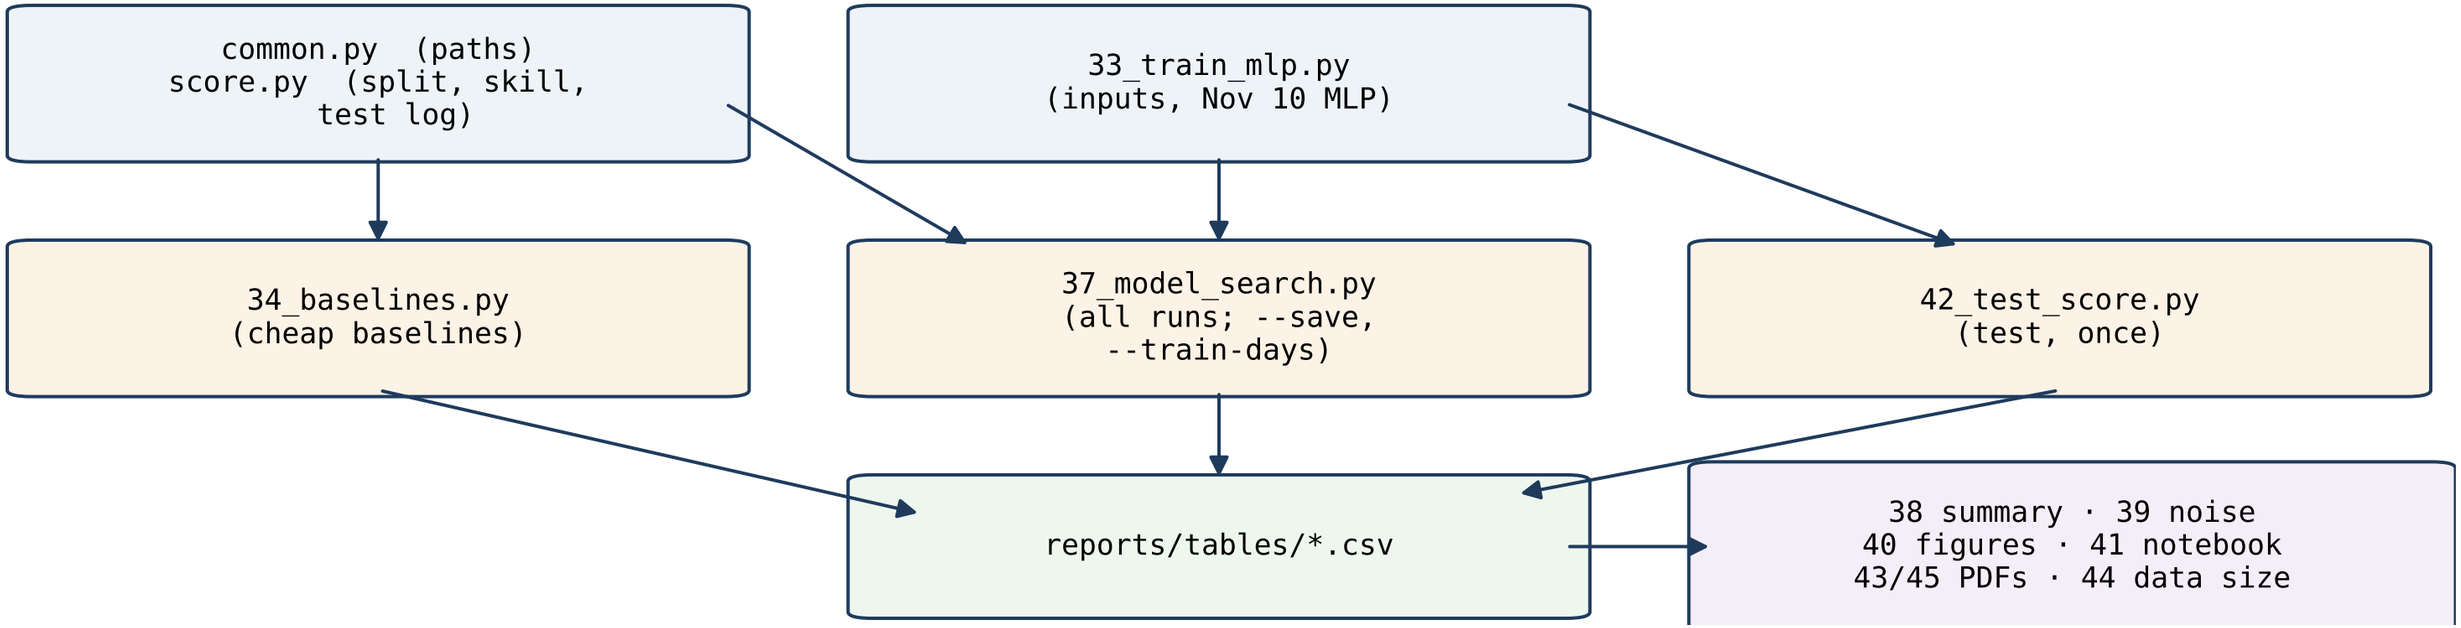
Test (read once): the e-bike gain holds (+0.013), the classic gain does not (+0.001); weekends score lower.



Test Oct 24–28, mean ± sd over seeds 0–2 (test_team_summary.csv, 42_test_score.py)

3. Code: verify it, change it

All in src/ of the project repo · <GitHub link: to be added (repo not decided yet)>



- Reproduce (from src/, Python env from requirements.txt): python 37_model_search.py --step 1 (or --run NAME --budget full --seed 0 1 2); python 38_search_report.py --finalists …; python 39_val_noise.py; python 44_data_size.py; python 42_test_score.py --dry (val checks only).
- Trace: notebooks/model_search_trace.ipynb (+ .html) prints every number of MODEL_SEARCH.md from its CSV with the script and commit; it checks the data SHA-1 and that the search never reads the test split.
- To change the model: add a config to GRIDS in 37_model_search.py; every run is logged to model_search_runs.csv. Data path: score.TEAM.